



Go from Data to Strategy ▶

ONLINE MASTER OF SCIENCE IN
BUSINESS ANALYTICSCarnegie Mellon University
Tepper School of Business[Go from Data to Strategy: Tepper School of Business](#)

Gradient Descent With Momentum from Scratch

by Jason Brownlee on [October 12, 2021](#) in [Optimization](#) 16

[f Share](#)[X Post](#)[in Share](#)

Gradient descent is an optimization algorithm that follows the negative gradient of an objective function in order to locate the minimum of the function.

A problem with gradient descent is that it can bounce around the search space on optimization problems that have large amounts of curvature or noisy gradients, and it can get stuck in flat spots in the search space that have no gradient.

Momentum is an extension to the gradient descent optimization algorithm that allows the search to build inertia in a direction in the search space and overcome the oscillations of noisy gradients and coast across flat spots of the search space.


In this tutorial, you will discover the gradient descent with momentum algorithm.

After completing this tutorial, you will know:

- Gradient descent is an optimization algorithm that uses the gradient of the objective function to navigate the search space.
- Gradient descent can be accelerated by using momentum from past updates to the search position.
- How to implement gradient descent optimization with momentum and develop an intuition for its behavior.

Kick-start your project with my new book [Optimization for Machine Learning](#), including *step-by-step tutorials* and the *Python source code* files for all examples.

Let's get started.



Gradient Descent With Momentum from Scratch
Photo by [Chris Barnes](#), some rights reserved.

Tutorial Overview

This tutorial is divided into three parts; they are:

1. Gradient Descent
2. Momentum
3. Gradient Descent With Momentum
 1. One-Dimensional Test Problem
 2. Gradient Descent Optimization
 3. Visualization of Gradient Descent Optimization
 4. Gradient Descent Optimization With Momentum
 5. Visualization of Gradient Descent Optimization With Momentum

Gradient Descent

Gradient descent is an optimization algorithm.

It is technically referred to as a first-order optimization algorithm as it explicitly makes use of the first-order derivative of the target objective function.



First-order methods rely on gradient information to help direct the search for a minimum ...

— Page 69, [Algorithms for Optimization](#), 2019.

The [first-order derivative](#), or simply the “*derivative*,” is the rate of change or slope of the target function at a specific point, e.g. for a specific input.

If the target function takes multiple input variables, it is referred to as a multivariate function and the input variables can be thought of as a vector. In turn, the derivative of a multivariate target function may also be taken as a vector and is referred to generally as the “[gradient](#).”

- **Gradient:** First-order derivative for a multivariate objective function.

The derivative or the gradient points in the direction of the steepest ascent of the target function for a specific input.

Gradient descent refers to a minimization optimization algorithm that follows the negative of the gradient downhill of the target function to locate the minimum of the function.

The gradient descent algorithm requires a target function that is being optimized and the derivative function for the objective function. The target function $f()$ returns a score for a given set of inputs, and the derivative function $f'()$ gives the derivative of the target function for a given set of inputs.

The gradient descent algorithm requires a starting point (x) in the problem, such as a randomly selected point in the input space.

The derivative is then calculated and a step is taken in the input space that is expected to result in a downhill movement in the target function, assuming we are minimizing the target function.

A downhill movement is made by first calculating how far to move in the input space, calculated as the step size (called *alpha* or the *learning rate*) multiplied by the gradient. This is then subtracted from the current point, ensuring we move against the gradient, or down the target function.

- $x = x - \text{step_size} * f'(x)$

The steeper the objective function at a given point, the larger the magnitude of the gradient and, in turn, the larger the step taken in the search space. The size of the step taken is scaled using a step size hyperparameter.

- **Step Size (*alpha*):** Hyperparameter that controls how far to move in the search space against the gradient each iteration of the algorithm, also called the learning rate.

If the step size is too small, the movement in the search space will be small and the search will take a long time. If the step size is too large, the search may bounce around the search space and skip over the optima.

Now that we are familiar with the gradient descent optimization algorithm, let's take a look at momentum.

Want to Get Started With Optimization Algorithms?

Take my free 7-day email crash course now (with sample code).

Click to sign-up and also get a free PDF Ebook version of the course.

Start your FREE Mini-Course now!

Momentum

Momentum is an extension to the gradient descent optimization algorithm, often referred to as **gradient descent with momentum**.

It is designed to accelerate the optimization process, e.g. decrease the number of function evaluations required to reach the optima, or to improve the capability of the optimization algorithm, e.g. result in a better final result.

A problem with the gradient descent algorithm is that the progression of the search can bounce around the search space based on the gradient. For example, the search may progress downhill towards the minima, but during this progression, it may move in another direction, even uphill, depending on the gradient of specific points (sets of parameters) encountered during the search.

This can slow down the progress of the search, especially for those optimization problems where the broader trend or shape of the search space is more useful than specific gradients along the way.

One approach to this problem is to add history to the parameter update equation based on the gradient encountered in the previous updates.

This change is based on the metaphor of momentum from physics where acceleration in a direction can be accumulated from past updates.



The name momentum derives from a physical analogy, in which the negative gradient is a force moving a particle through parameter space, according to Newton's laws of motion.

— Page 296, [Deep Learning](#), 2016.

Momentum involves adding an additional hyperparameter that controls the amount of history (momentum) to include in the update equation, i.e. the step to a new point in the search space. The value for the hyperparameter is defined in the range 0.0 to 1.0 and often has a value close to 1.0, such as 0.8, 0.9, or 0.99. A momentum of 0.0 is the same as gradient descent without momentum.

First, let's break the gradient descent update equation down into two parts: the calculation of the change to the position and the update of the old position to the new position.

The change in the parameters is calculated as the gradient for the point scaled by the step size.

- $\text{change_x} = \text{step_size} * f'(x)$

The new position is calculated by simply subtracting the change from the current point

- $x = x - \text{change_x}$

Momentum involves maintaining the change in the position and using it in the subsequent calculation of the change in position.

If we think of updates over time, then the update at the current iteration or time (t) will add the change used at the previous time (t-1) weighted by the momentum hyperparameter, as follows:

- $\text{change_x}(t) = \text{step_size} * f'(x(t-1)) + \text{momentum} * \text{change_x}(t-1)$

The update to the position is then performed as before.

- $x(t) = x(t-1) - \text{change_x}(t)$

The change in the position accumulates magnitude and direction of changes over the iterations of the search, proportional to the size of the momentum hyperparameter.

For example, a large momentum (e.g. 0.9) will mean that the update is strongly influenced by the previous update, whereas a modest momentum (0.2) will mean very little influence.



The momentum algorithm accumulates an exponentially decaying moving average of past gradients and continues to move in their direction.

— Page 296, [Deep Learning](#), 2016.

Momentum has the effect of dampening down the change in the gradient and, in turn, the step size with each new point in the search space.



Momentum can increase speed when the cost surface is highly nonspherical because it damps the size of the steps along directions of high curvature thus yielding a larger effective learning rate along the directions of low curvature.

— Page 21, [Neural Networks: Tricks of the Trade](#), 2012.

Momentum is most useful in optimization problems where the objective function has a large amount of curvature (e.g. changes a lot), meaning that the gradient may change a lot over relatively small regions of the search space.



The method of momentum is designed to accelerate learning, especially in the face of high curvature, small but consistent gradients, or noisy gradients.

— Page 296, [Deep Learning](#), 2016.

It is also helpful when the gradient is estimated, such as from a simulation, and may be noisy, e.g. when the gradient has a high variance.

Finally, momentum is helpful when the search space is flat or nearly flat, e.g. zero gradient. The momentum allows the search to progress in the same direction as before the flat spot and helpfully cross the flat region.

Now that we are familiar with what momentum is, let's look at a worked example.

Gradient Descent with Momentum

In this section, we will first implement the gradient descent optimization algorithm, then update it to use momentum and compare results.

One-Dimensional Test Problem

First, let's define an optimization function.

We will use a simple one-dimensional function that squares the input and defines the range of valid inputs from -1.0 to 1.0.

The `objective()` function below implements this function.

```
1 # objective function
2 def objective(x):
3     return x**2.0
```

We can then sample all inputs in the range and calculate the objective function value for each.

```
1 ...
2 # define range for input
3 r_min, r_max = -1.0, 1.0
4 # sample input range uniformly at 0.1 increments
5 inputs = arange(r_min, r_max+0.1, 0.1)
6 # compute targets
7 results = objective(inputs)
```

Finally, we can create a line plot of the inputs (x-axis) versus the objective function values (y-axis) to get an intuition for the shape of the objective function that we will be searching.

```
1 ...
2 # create a line plot of input vs result
3 pyplot.plot(inputs, results)
4 # show the plot
5 pyplot.show()
```

The example below ties this together and provides an example of plotting the one-dimensional test function.

```
1 # plot of simple function
2 from numpy import arange
3 from matplotlib import pyplot
4
5 # objective function
6 def objective(x):
7     return x**2.0
8
9 # define range for input
10 r_min, r_max = -1.0, 1.0
11 # sample input range uniformly at 0.1 increments
12 inputs = arange(r_min, r_max+0.1, 0.1)
13 # compute targets
14 results = objective(inputs)
15 # create a line plot of input vs result
16 pyplot.plot(inputs, results)
17 # show the plot
18 pyplot.show()
```

Running the example creates a line plot of the inputs to the function (x-axis) and the calculated output of the function (y-axis).

We can see the familiar U-shape called a parabola.

Line Plot of Simple One Dimensional Function

Gradient Descent Optimization

Next, we can apply the gradient descent algorithm to the problem.

First, we need a function that calculates the derivative for the objective function.

The derivative of x^2 is $x * 2$ and the `derivative()` function implements this below.

```
1 # derivative of objective function
2 def derivative(x):
3     return x * 2.0
```

We can define a function that implements the gradient descent optimization algorithm.

The procedure involves starting with a randomly selected point in the search space, then calculating the gradient, updating the position in the search space, evaluating the new position, and reporting the progress. This process is then repeated for a fixed number of iterations. The final point and its evaluation are then returned from the function.

The function `gradient_descent()` below implements this and takes the name of the objective and gradient functions as well as the bounds on the inputs to the objective function, number of iterations, and step size, then returns the solution and its evaluation at the end of the search.

```
1 # gradient descent algorithm
2 def gradient_descent(objective, derivative, bounds, n_iter, step_size):
3     # generate an initial point
4     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
5     # run the gradient descent
6     for i in range(n_iter):
7         # calculate gradient
8         gradient = derivative(solution)
9         # take a step
10        solution = solution - step_size * gradient
11        # evaluate candidate point
12        solution_eval = objective(solution)
13        # report progress
14        print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
15    return [solution, solution_eval]
```

We can then define the bounds of the objective function, the step size, and the number of iterations for the algorithm.

We will use a step size of 0.1 and 30 iterations, both found after a little experimentation.

The seed for the pseudorandom number generator is fixed so that we always get the same sequence of random numbers, and in this case, it ensures that we get the same starting point for the search each time the code is run (e.g. something interesting far from the optima).

```

1 ...
2 # seed the pseudo random number generator
3 seed(4)
4 # define range for input
5 bounds = asarray([[ -1.0, 1.0]])
6 # define the total iterations
7 n_iter = 30
8 # define the maximum step size
9 step_size = 0.1
10 # perform the gradient descent search
11 best, score = gradient_descent(objective, derivative, bounds, n_iter, step_size)

```

Tying this together, the complete example of applying grid search to our one-dimensional test function is listed below.

```

1 # example of gradient descent for a one-dimensional function
2 from numpy import asarray
3 from numpy.random import rand
4 from numpy.random import seed
5
6 # objective function
7 def objective(x):
8     return x**2.0
9
10 # derivative of objective function
11 def derivative(x):
12     return x * 2.0
13
14 # gradient descent algorithm
15 def gradient_descent(objective, derivative, bounds, n_iter, step_size):
16     # generate an initial point
17     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
18     # run the gradient descent
19     for i in range(n_iter):
20         # calculate gradient
21         gradient = derivative(solution)
22         # take a step
23         solution = solution - step_size * gradient
24         # evaluate candidate point
25         solution_eval = objective(solution)
26         # report progress
27         print('%d f(%s) = %.5f' % (i, solution, solution_eval))
28     return [solution, solution_eval]
29
30 # seed the pseudo random number generator
31 seed(4)
32 # define range for input
33 bounds = asarray([[ -1.0, 1.0]])
34 # define the total iterations
35 n_iter = 30
36 # define the step size
37 step_size = 0.1
38 # perform the gradient descent search
39 best, score = gradient_descent(objective, derivative, bounds, n_iter, step_size)
40 print('Done!')
41 print('f(%s) = %f' % (best, score))

```

Running the example starts with a random point in the search space, then applies the gradient descent algorithm, reporting performance along the way.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

In this case, we can see that the algorithm finds a good solution after about 27 iterations, with a function evaluation of about 0.0.

Note the optima for this function is at $f(0.0) = 0.0$.

We would expect that gradient descent with momentum will accelerate the optimization procedure and find a similarly evaluated solution in fewer iterations.

```

1 >0 f([0.74724774]) = 0.55838
2 >1 f([0.59779819]) = 0.35736
3 >2 f([0.47823856]) = 0.22871
4 >3 f([0.38259084]) = 0.14638
5 >4 f([0.30607268]) = 0.09368
6 >5 f([0.24485814]) = 0.05996
7 >6 f([0.19588651]) = 0.03837
8 >7 f([0.15670921]) = 0.02456
9 >8 f([0.12536737]) = 0.01572
10 >9 f([0.10029389]) = 0.01006
11 >10 f([0.08023512]) = 0.00644
12 >11 f([0.06418809]) = 0.00412
13 >12 f([0.05135047]) = 0.00264
14 >13 f([0.04108038]) = 0.00169
15 >14 f([0.0328643]) = 0.00108
16 >15 f([0.02629144]) = 0.00069
17 >16 f([0.02103315]) = 0.00044
18 >17 f([0.01682652]) = 0.00028
19 >18 f([0.01346122]) = 0.00018
20 >19 f([0.01076897]) = 0.00012

```

```

21 >20 f([0.00861518]) = 0.00007
22 >21 f([0.00689214]) = 0.00005
23 >22 f([0.00551372]) = 0.00003
24 >23 f([0.00441097]) = 0.00002
25 >24 f([0.00352878]) = 0.00001
26 >25 f([0.00282302]) = 0.00001
27 >26 f([0.00225842]) = 0.00001
28 >27 f([0.00180673]) = 0.00000
29 >28 f([0.00144539]) = 0.00000
30 >29 f([0.00115631]) = 0.00000
31 Done!
32 f([0.00115631]) = 0.000001

```

Visualization of Gradient Descent Optimization

Next, we can visualize the progress of the search on a plot of the target function.

First, we can update the `gradient_descent()` function to store all solutions and their score found during the optimization as lists and return them at the end of the search instead of the best solution found.

```

1 # gradient descent algorithm
2 def gradient_descent(objective, derivative, bounds, n_iter, step_size):
3     # track all solutions
4     solutions, scores = list(), list()
5     # generate an initial point
6     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
7     # run the gradient descent
8     for i in range(n_iter):
9         # calculate gradient
10        gradient = derivative(solution)
11        # take a step
12        solution = solution - step_size * gradient
13        # evaluate candidate point
14        solution_eval = objective(solution)
15        # store solution
16        solutions.append(solution)
17        scores.append(solution_eval)
18        # report progress
19        print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
20    return [solutions, scores]

```

The function can be called and we can get the lists of the solutions and the scores found during the search.

```

1 ...
2 # perform the gradient descent search
3 solutions, scores = gradient_descent(objective, derivative, bounds, n_iter, step_size)

```

We can create a line plot of the objective function, as before.

```

1 ...
2 # sample input range uniformly at 0.1 increments
3 inputs = arange(bounds[0,0], bounds[0,1]+0.1, 0.1)
4 # compute targets
5 results = objective(inputs)
6 # create a line plot of input vs result
7 pyplot.plot(inputs, results)

```

Finally, we can plot each solution found as a red dot and connect the dots with a line so we can see how the search moved downhill.

```

1 ...
2 # plot the solutions found
3 pyplot.plot(solutions, scores, '.-', color='red')

```

Tying this all together, the complete example of plotting the result of the gradient descent search on the one-dimensional test function is listed below.

```

1 # example of plotting a gradient descent search on a one-dimensional function
2 from numpy import asarray
3 from numpy import arange
4 from numpy.random import rand
5 from numpy.random import seed
6 from matplotlib import pyplot
7
8 # objective function
9 def objective(x):
10    return x**2.0
11
12 # derivative of objective function
13 def derivative(x):
14    return x * 2.0
15
16 # gradient descent algorithm
17 def gradient_descent(objective, derivative, bounds, n_iter, step_size):
18     # track all solutions
19     solutions, scores = list(), list()
20     # generate an initial point
21     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
22     # run the gradient descent
23     for i in range(n_iter):
24         # calculate gradient
25         gradient = derivative(solution)
26         # take a step
27         solution = solution - step_size * gradient
28         # evaluate candidate point
29         solution_eval = objective(solution)
30         # store solution
31         solutions.append(solution)

```



```

32     scores.append(solution_eval)
33     # report progress
34     print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
35     return [solutions, scores]
36
37 # seed the pseudo random number generator
38 seed(4)
39 # define range for input
40 bounds = asarray([[-1.0, 1.0]])
41 # define the total iterations
42 n_iter = 30
43 # define the step size
44 step_size = 0.1
45 # perform the gradient descent search
46 solutions, scores = gradient_descent(objective, derivative, bounds, n_iter, step_size)
47 # sample input range uniformly at 0.1 increments
48 inputs = arange(bounds[0,0], bounds[0,1]+0.1, 0.1)
49 # compute targets
50 results = objective(inputs)
51 # create a line plot of input vs result
52 pyplot.plot(inputs, results)
53 # plot the solutions found
54 pyplot.plot(solutions, scores, '.-', color='red')
55 # show the plot
56 pyplot.show()

```

Running the example performs the gradient descent search on the objective function as before, except in this case, each point found during the search is plotted.

Note: Your [results may vary](#) given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

In this case, we can see that the search started more than halfway up the right part of the function and stepped downhill to the bottom of the basin.

We can see that in the parts of the objective function with the larger curve, the derivative (gradient) is larger, and in turn, larger steps are taken. Similarly, the gradient is smaller as we get closer to the optima, and in turn, smaller steps are taken.

This highlights that the step size is used as a scale factor on the magnitude of the gradient (curvature) of the objective function.

Plot of the Progress of Gradient Descent on a One Dimensional Objective Function

Gradient Descent Optimization With Momentum

Next, we can update the gradient descent optimization algorithm to use momentum.

This can be achieved by updating the `gradient_descent()` function to take a “*momentum*” argument that defines the amount of momentum used during the search.

The change made to the solution must be remembered from the previous iteration of the loop, with an initial value of 0.0.

```
1 ...
2 # keep track of the change
3 change = 0.0
```

We can then break the update procedure down into first calculating the gradient, then calculating the change to the solution, calculating the position of the new solution, then saving the change for the next iteration.

```
1 ...
2 # calculate gradient
3 gradient = derivative(solution)
4 # calculate update
5 new_change = step_size * gradient + momentum * change
6 # take a step
7 solution = solution - new_change
8 # save the change
9 change = new_change
```

The updated version of the `gradient_descent()` function with these changes is listed below.

```
1 # gradient descent algorithm
2 def gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum):
3     # generate an initial point
4     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
5     # keep track of the change
6     change = 0.0
7     # run the gradient descent
8     for i in range(n_iter):
9         # calculate gradient
10        gradient = derivative(solution)
11        # calculate update
12        new_change = step_size * gradient + momentum * change
13        # take a step
14        solution = solution - new_change
15        # save the change
16        change = new_change
17        # evaluate candidate point
18        solution_eval = objective(solution)
19        # report progress
20        print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
21    return [solution, solution_eval]
```

We can then choose a momentum value and pass it to the `gradient_descent()` function.

After a little trial and error, a momentum value of 0.3 was found to be effective on this problem, given the fixed step size of 0.1.

```
1 ...
2 # define momentum
3 momentum = 0.3
4 # perform the gradient descent search with momentum
5 best, score = gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum)
```

Tying this together, the complete example of gradient descent optimization with momentum is listed below.

```
1 # example of gradient descent with momentum for a one-dimensional function
2 from numpy import asarray
3 from numpy.random import rand
4 from numpy.random import seed
5
6 # objective function
7 def objective(x):
8     return x**2.0
9
10 # derivative of objective function
11 def derivative(x):
12     return x * 2.0
13
14 # gradient descent algorithm
15 def gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum):
16     # generate an initial point
17     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
18     # keep track of the change
19     change = 0.0
20     # run the gradient descent
21     for i in range(n_iter):
22         # calculate gradient
23         gradient = derivative(solution)
24         # calculate update
25         new_change = step_size * gradient + momentum * change
26         # take a step
27         solution = solution - new_change
28         # save the change
29         change = new_change
30         # evaluate candidate point
31         solution_eval = objective(solution)
32         # report progress
33         print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
34     return [solution, solution_eval]
35
```

```

36 # seed the pseudo random number generator
37 seed(4)
38 # define range for input
39 bounds = asarray([[-1.0, 1.0]])
40 # define the total iterations
41 n_iter = 30
42 # define the step size
43 step_size = 0.1
44 # define momentum
45 momentum = 0.3
46 # perform the gradient descent search with momentum
47 best, score = gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum)
48 print('Done!')
49 print('f(%s) = %f' % (best, score))

```

Running the example starts with a random point in the search space, then applies the gradient descent algorithm with momentum, reporting performance along the way.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

In this case, we can see that the algorithm finds a good solution after about 13 iterations, with a function evaluation of about 0.0.

As expected, this is faster (fewer iterations) than gradient descent without momentum, using the same starting point and step size that took 27 iterations.

```

1 >0 f([0.74724774]) = 0.55838
2 >1 f([0.54175461]) = 0.29350
3 >2 f([0.37175575]) = 0.13820
4 >3 f([0.24640494]) = 0.06072
5 >4 f([0.15951871]) = 0.02545
6 >5 f([0.1015491]) = 0.01031
7 >6 f([0.0638484]) = 0.00408
8 >7 f([0.03976851]) = 0.00158
9 >8 f([0.02459084]) = 0.00060
10 >9 f([0.01511937]) = 0.00023
11 >10 f([0.00925406]) = 0.00009
12 >11 f([0.00564365]) = 0.00003
13 >12 f([0.0034318]) = 0.00001
14 >13 f([0.00208188]) = 0.00000
15 >14 f([0.00126053]) = 0.00000
16 >15 f([0.00076202]) = 0.00000
17 >16 f([0.00046006]) = 0.00000
18 >17 f([0.00027746]) = 0.00000
19 >18 f([0.00016719]) = 0.00000
20 >19 f([0.00010067]) = 0.00000
21 >20 f([6.05804744e-05]) = 0.00000
22 >21 f([3.64373635e-05]) = 0.00000
23 >22 f([2.19069576e-05]) = 0.00000
24 >23 f([1.31664443e-05]) = 0.00000
25 >24 f([7.91100141e-06]) = 0.00000
26 >25 f([4.75216828e-06]) = 0.00000
27 >26 f([2.85408468e-06]) = 0.00000
28 >27 f([1.71384267e-06]) = 0.00000
29 >28 f([1.02900153e-06]) = 0.00000
30 >29 f([6.17748881e-07]) = 0.00000
31 Done!
32 f([6.17748881e-07]) = 0.00000

```

Visualization of Gradient Descent Optimization With Momentum

Finally, we can visualize the progress of the gradient descent optimization algorithm with momentum.

The complete example is listed below.

```

1 # example of plotting gradient descent with momentum for a one-dimensional function
2 from numpy import asarray
3 from numpy import arange
4 from numpy.random import rand
5 from numpy.random import seed
6 from matplotlib import pyplot
7
8 # objective function
9 def objective(x):
10     return x**2.0
11
12 # derivative of objective function
13 def derivative(x):
14     return x * 2.0
15
16 # gradient descent algorithm
17 def gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum):
18     # track all solutions
19     solutions, scores = list(), list()
20     # generate an initial point
21     solution = bounds[:, 0] + rand(len(bounds)) * (bounds[:, 1] - bounds[:, 0])
22     # keep track of the change
23     change = 0.0
24     # run the gradient descent
25     for i in range(n_iter):
26         # calculate gradient
27         gradient = derivative(solution)
28         # calculate update
29         new_change = step_size * gradient + momentum * change
30         # take a step

```

```

31     solution = solution - new_change
32     # save the change
33     change = new_change
34     # evaluate candidate point
35     solution_eval = objective(solution)
36     # store solution
37     solutions.append(solution)
38     scores.append(solution_eval)
39     # report progress
40     print('>%d f(%s) = %.5f' % (i, solution, solution_eval))
41     return [solutions, scores]
42
43 # seed the pseudo random number generator
44 seed(4)
45 # define range for input
46 bounds = asarray([[-1.0, 1.0]])
47 # define the total iterations
48 n_iter = 30
49 # define the step size
50 step_size = 0.1
51 # define momentum
52 momentum = 0.3
53 # perform the gradient descent search with momentum
54 solutions, scores = gradient_descent(objective, derivative, bounds, n_iter, step_size, momentum)
55 # sample input range uniformly at 0.1 increments
56 inputs = arange(bounds[0,0], bounds[0,1]+0.1, 0.1)
57 # compute targets
58 results = objective(inputs)
59 # create a line plot of input vs result
60 pyplot.plot(inputs, results)
61 # plot the solutions found
62 pyplot.plot(solutions, scores, '.-', color='red')
63 # show the plot
64 pyplot.show()

```

Running the example performs the gradient descent search with momentum on the objective function as before, except in this case, each point found during the search is plotted.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

In this case, if we compare the plot to the plot created previously for the performance of gradient descent (without momentum), we can see that the search indeed reaches the optima in fewer steps, noted with fewer distinct red dots on the path to the bottom of the basin.

Plot of the Progress of Gradient Descent With Momentum on a One Dimensional Objective Function

As an extension, try different values for momentum, such as 0.8, and review the resulting plot.
Let me know what you discover in the comments below.

Further Reading

This section provides more resources on the topic if you are looking to go deeper.

Books

- [Algorithms for Optimization](#), 2019.
- [Deep Learning](#), 2016.
- [Neural Smithing: Supervised Learning in Feedforward Artificial Neural Networks](#), 1999.
- [Neural Networks for Pattern Recognition](#), 1996.
- [Neural Networks: Tricks of the Trade](#), 2012.

APIs

- [numpy.random.rand API](#).
- [numpy.asarray API](#).
- [Matplotlib API](#).

Articles

- [Gradient descent, Wikipedia](#).
- [Stochastic gradient descent, Wikipedia](#).
- [Gradient, Wikipedia](#).
- [Derivative, Wikipedia](#).
- [Differentiable function, Wikipedia](#).

Summary

In this tutorial, you discovered the gradient descent with momentum algorithm.

Specifically, you learned:

- Gradient descent is an optimization algorithm that uses the gradient of the objective function to navigate the search space.
- Gradient descent can be accelerated by using momentum from past updates to the search position.
- How to implement gradient descent optimization with momentum and develop an intuition for its behavior.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

Get a Handle on Modern Optimization Algorithms!

Develop Your Understanding of Optimization

...with just a few lines of python code

Discover how in my new Ebook:
[Optimization for Machine Learning](#)







It provides **self-study tutorials** with **full working code** on:
Gradient Descent, Genetic Algorithms, Hill Climbing, Curve Fitting, RMSProp, Adam, and much more...

Bring Modern Optimization Algorithms to Your Machine Learning Projects

SEE WHAT'S INSIDE

[f Share](#) [X Post](#) [in Share](#)

More On This Topic



Gradient Descent With Adadelta
from Scratch

Gradient Descent Optimization
With AdaMax From Scratch

About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

< [Weight Initialization for Deep Learning Neural Networks](#)

[Function Optimization With SciPy](#) >

16 Responses to *Gradient Descent With Momentum from Scratch*

Hugh Brown February 5, 2021 at 2:01 pm #

REPLY ↩

site that has this content:

<https://neoshare.net/machine-learning/gradient-descent-with-momentum-from-scratch/>

Jason Brownlee February 5, 2021 at 5:00 pm #

REPLY ↩

Stolen directly from my site!!!

Don't worry, google punishes these copy-cats (rip-off artists) with terrible search rankings...

Beliaivsky February 6, 2021 at 8:49 am #

REPLY ↩

Could you next present the case of minimizing a function of several variables? Thanks for your post.

Jason Brownlee February 6, 2021 at 2:22 pm #

REPLY ↩

Sure, see this example:

<https://machinelearningmastery.com/adam-optimization-from-scratch/>

fabou February 7, 2021 at 10:13 pm #

REPLY ↩

Hi Jason, nice post as usual.

What if you do not have the derivative of your function ? What would you recommend ?

Choosing a good momentum value seems to be an optimization problem as well. Could momentum algorithm be used in a recursive way?

Jason Brownlee February 8, 2021 at 7:02 am #

REPLY ↩

You can use a pattern search like nelder mead:
<https://machinelearningmastery.com/how-to-use-nelder-mead-optimization-in-python/>
 Or a stochastic search like stochastic hill climbing:
<https://machinelearningmastery.com/stochastic-hill-climbing-in-python-from-scratch/>

Yes, momentum is yet another hyperparameter to tune.

Vince October 13, 2021 at 1:35 am #

REPLY ↩

Thanks for the nice example. I am currently using finite difference to approximate my gradient in a simulation optimization. Let's see if the momentum term speeds up my optimizaton 😊

Adrian Tam October 13, 2021 at 7:33 am #

REPLY ↩

Please try. My impression is that momentum may not speed up, but allows you to get out of premature convergence in case you have local optima.

JAIKISHAN February 12, 2021 at 3:54 pm #

REPLY ↩

Thanks Jason for the vivid explanation
 Are there any other hyper parameters in the optimizer function other than the Learning rate and momentum used in regular applications in the industry??

Jason Brownlee February 13, 2021 at 5:57 am #

REPLY ↩

You're welcome.
 No, not really.

Nishant Gaurav August 7, 2021 at 10:49 am #

REPLY ↩

Hi Jason, I have been following your work for a very long. I have also bought few books of yours. Jason, will it be possible for you to showcase these with a dataset having just input features of X and output features of Y. It can help us in relating with the concepts in a better way.

Jason Brownlee August 8, 2021 at 5:07 am #

REPLY ↩

Thanks for the suggestion!

Denis February 5, 2022 at 10:36 pm #

REPLY ↩

Hi Jason
 Great article, thanks for this.
 One question. There is a formula in the beginning of the article explaining how momentum works. See below

$$\text{change_x}(t) = \text{step_size} * f'(x(t-1)) + \text{momentum} * \text{change_x}(t-1)$$

 Shouldn't it be like this

$$\text{change_x}(t) = \text{step_size} * f'(x(t)) + \text{momentum} * \text{change_x}(t-1)$$

 If not, I little bit confused then)

James Carmichael February 17, 2022 at 1:34 pm #

REPLY ↩

Hi Denis...You are correct. This has been noted.

Tom August 11, 2023 at 6:50 am <#>

REPLY ↩

Doesn't this effectively sum the gradients and so causing exploding gradients?

James Carmichael August 11, 2023 at 8:17 am <#>

REPLY ↩


Hi Tom...The following resource explains the benefits of gradient descent with momentum from another perspective.

<https://distill.pub/2017/momentum/>

Leave a Reply

 Name (required) Email (will not be published) (required)

SUBMIT COMMENT



Welcome!

I'm *Jason Brownlee* PhD

and I **help developers** get results with **machine learning**.

[Read more](#)

Never miss a tutorial:



Picked for you:



[Simple Genetic Algorithm From Scratch in Python](#)



[Curve Fitting With Python](#)



[A Gentle Introduction to Particle Swarm Optimization](#)



[Optimization for Machine Learning Crash Course](#)



Code Adam Optimization Algorithm From Scratch

Loving the Tutorials?

The [Optimization for Machine Learning](#) EBook
is where you'll find the ***Really Good*** stuff.

[>> SEE WHAT'S INSIDE](#)

Machine Learning Mastery is part of Guiding Tech Media, a leading digital media publisher focused on helping people figure out technology. [Visit our corporate website](#) to learn more about our mission and team.



[PRIVACY](#) | [DISCLAIMER](#) | [TERMS](#) | [CONTACT](#) | [SITEMAP](#)

© 2025 Guiding Tech Media All Rights Reserved
